

SECURITY + AI

SECURITY VAULT

Five underground security & AI tools most people never find — with a ready-to-paste prompt for each.

TUNED FOR MACOS • zsh

JUNE 2026 • 5 PICKS • EN / AR

> BitNet > context-mode > cmux > Scrapy > mcp-scan



Welcome to the Vault.

Every week you get one short PDF: a handful of genuinely useful security & AI tools — the under-surfaced ones, not the stuff already on every feed. Each pick comes with what it does, why it matters for you, the real GitHub link, and a prompt you paste straight into your AI to actually use it.

► HOW TO USE EACH PICK

- 1. Read the hook — decide in 5 seconds if it's for you.
- 2. Open the GitHub link.
- 3. Copy the **>_ DROP THIS INTO YOUR AI** prompt into Claude, ChatGPT, or Codex.
- 4. Let it walk you through setup & first use. No guesswork.

► YOUR BUILD • MACOS Commands are written for zsh — install prerequisites with Homebrew (brew install ...).

كل أسبوع ملف واحد مختصر: أدوات أمن وذكاء اصطناعي مفيدة فعلاً — السرية منها، مش اللي تشوفها بكل مكان. كل أداة معها شرح، وليش تهتمك أنت، ورابط GitHub الحقيقي، وبرومبت جاهز تنسخه في Claude أو ChatGPT أو Codex وتبدأ على طول.

IN THIS VAULT

01	BitNet	LOCAL AI • NO GPU
02	context-mode	AI SESSION INFRA
03	cmux	MULTI-AGENT TERMINAL
04	Scrapy	AUTOMATION BACKBONE
05	mcp-scan	MCP SECURITY

PRIVACY / LOCAL AI • MICROSOFT

BitNet

01

★ ~22k • Microsoft official • verdict [SHARE](#)

Run real LLMs on any CPU — no GPU, no cloud. Microsoft's 1-bit models are light enough to run interactively on a plain laptop, and your AI never leaves your machine.

ذكاء اصطناعي بدون GPU — مايكروسوفت تطلق نماذج تشتغل على أي جهاز. شغل نماذج لغوية محليًا على المعالج فقط، بدون كرت شاشة وبدون سحابة، وخصوصيتك محفوظة.

WHAT IT IS

bitnet.cpp is Microsoft's official inference framework for 1-bit / 1.58-bit LLMs — weights stored as just -1, 0, or +1 instead of 16–32 bits. That extreme compression means a 2B model runs fast and snappy on plain x86 CPU (2–6× faster than standard llama.cpp), and even 100B-class models can run on a single CPU at reading speed.

WHY IT MATTERS FOR YOU

No \$2/hr cloud GPU, no VRAM panic — clone it and run a local chatbot today. Perfect for private experiments, offline tools, or anything where data can't leave the device. (The fast GPU path is CUDA-only, so on AMD you stay on the CPU path — which is the point here anyway.)

```
repo > github.com/microsoft/BitNet
```

```
>_ DROP THIS INTO YOUR AI • macOS / zsh
```

Set up Microsoft BitNet (github.com/microsoft/BitNet) to run a 1-bit LLM on my CPU with no GPU. Give me the exact Windows steps: clone the repo, build bitnet.cpp, download the official 2B model from Hugging Face, and run an interactive chat. Flag any prerequisite I'm missing before I start.

AI TOOLS / CLAUDE CODE • MCP

02

context-mode

★ ~2.5k • #1 on Hacker News • verdict BOTH

Stop losing your AI coding session every 30 minutes. context-mode sandboxes tool output and survives context compaction — turning 30-minute sessions into 3-hour ones.

كل ٣٠ دقيقة تشتغل مع كلود والسياق يموت؟ context-mode يخليك تشتغل ٣ ساعات بدون ما تعيد من الأول — طبقة البنية التحتية اللي المحترفون يستخدمونها فعلاً.

WHAT IT IS

An MCP server + plugin that stops raw tool output from flooding your context window — a 56 KB Playwright snapshot becomes 299 bytes. It also persists every edit, task, and decision to a local SQLite DB and rebuilds your working state automatically when a session compacts. Works with Claude Code, Cursor, Gemini CLI, Codex, Copilot, Zed.

WHY IT MATTERS FOR YOU

If you run real multi-step AI sessions, you hit the context wall constantly. This is the fix — one install, hooks in automatically, no CLAUDE.md edits. Genuinely under-surfaced: ~2.5k stars despite #1 on HN and 240k+ users.

```
repo > github.com/mksglu/context-mode
```

```
>_ DROP THIS INTO YOUR AI • macOS / zsh
```

Install context-mode (github.com/mksglu/context-mode) in my Claude Code setup so sessions stop dying at the context limit. Walk me through adding the plugin marketplace, enabling it, and confirming the PreToolUse/PreCompact hooks are active. Then show me how to read the savings with `ctx_stats`.

AI TOOLS / TERMINAL • macOS

03

cmux

★ ~20k • native macOS • Ghostty • verdict BOTH

Running 10 AI agents and babysitting every terminal? cmux gives each Claude / Codex / Gemini session a notification ring, so you know exactly which one needs you.

بدل ما تراقب ١٠ نوافذ ترمينال وتنتظر إشعارات cmux — Claude يحط حلقة إشعار لكل جلسة تعرف منها أيها يحتاجك. مبني من مهندسين يشغلون agents فعلاً.

WHAT IT IS

A native macOS terminal (Swift + libghostty) with vertical sidebar tabs that show each workspace's git branch, ports, and latest agent notification. A built-in scriptable browser pane lets Claude Code click, fill forms, and run JS without leaving the terminal. cmux claude-teams spawns teammate mode as native splits — no tmux.

WHY IT MATTERS FOR YOU

The moment you run more than two agents in parallel, generic macOS notifications become useless — they're all identical. cmux solves exactly that. It reads your existing Ghostty config, so zero reconfiguration. (macOS only.)

```
repo > github.com/manaflo-ai/cmux
```

```
>_ DROP THIS INTO YOUR AI • macOS / zsh
```

I run multiple Claude Code / Codex agents at once on macOS and lose track of which needs me. Set up cmux (github.com/manaflo-ai/cmux): install via Homebrew, wire the cmux notify ring into my agent hooks, and explain how cmux claude-teams works for parallel teammate sessions.

AUTOMATION / DATA • PYTHON

Scrapy

04

★ ~55k • maintained by Zyte • verdict BOTH

When a scraping script stops being enough, Scrapy is the production backbone — spiders, pipelines, throttling, and scheduling built in. The data layer behind serious AI datasets and market intel.

Scrapy مش مجرد أداة تسحب بيانات — هو الإطار اللي تبني فوقه أي مشروع automation جاد: من جمع leads، لمراقبة الأسعار، لبناء datasets لتدريب الـ AI. لو لسه تكتب requests يدوية، أنت خسران وقت كل يوم.

WHAT IT IS

A battle-tested Python framework that extracts structured data from any site with a clean spider architecture — concurrency, middleware, item pipelines, and export to JSON/CSV/XML out of the box, from the CLI or as a library. Pair it with rotating proxies + a scheduler and it covers ~80% of data-collection jobs without paying for SaaS.

WHY IT MATTERS FOR YOU

Any content aggregator, price monitor, lead-gen pipeline, or AI training set eventually outgrows one-off scripts. Positioned right — not "a scraper" but "the automation backbone" — this is foundational, high-leverage infrastructure.

```
repo > github.com/scrapy/scrapy
```

```
>_ DROP THIS INTO YOUR AI • macOS / zsh
```

I'm outgrowing one-off requests scripts. Scaffold a production Scrapy project (github.com/scrapy/scrapy) for the site I'll name: a spider, an item pipeline that exports JSON, polite throttling + auto-throttle, and how to schedule recurring runs. Explain each piece as you build it.

SECURITY / MCP • INVARIANT → SNYK

05

mcp-scan

★ ~2.6k • coined "Tool Poisoning" • acq. by Snyk • verdict BOTH

That MCP server you just installed can read your SSH keys — and your agent will never tell you. mcp-scan catches "Tool Poisoning": hidden instructions inside tool descriptions the model silently obeys.

أي MCP server ترغّبه على Claude أو Cursor ممكن يكون مفعّخ — يقرأ مفاتيحك بصمت وال agent يرجّعلك إجابة عادية. mcp-scan يكشف هجمات Tool Poisoning قبل ما تثق بأي سيرفر، ومكتشفوه أثبتوا الهجوم على سيرفرات WhatsApp و GitHub نفسها.

WHAT IT IS

A security scanner for the Model Context Protocol — the plumbing that lets agents call external tools. It auto-finds your MCP configs (Claude Desktop, Cursor, VS Code...), pulls each server's tool descriptions, and flags tool poisoning, tool shadowing, rug pulls, and cross-origin escalations. Invariant Labs named the attack class and was acquired by Snyk in 2025.

WHY IT MATTERS FOR YOU

Every MCP server you add is an unvetted trust boundary. A 30-second scan over your config catches a poisoned tool before it ever reaches a session. The security bookend to picks 02 & 03 — those run your agents; this keeps them from getting hijacked.

```
repo > github.com/invariantlabs-ai/mcp-scan
```

```
>_ DROP THIS INTO YOUR AI • macOS / zsh
```

```
Audit my MCP servers for security before I trust them. Walk me through using mcp-scan (github.com/invariantlabs-ai/mcp-scan) to scan my Claude Desktop / Cursor config for tool poisoning and rug pulls — the exact command to run, what each finding type means, and what to do if it flags one of my servers.
```

That's the Vault for this week.

Five tools, five prompts. Try even one of them this week — that's the whole point. Real tools you actually run, not another feed to scroll.

Next week: more under-surfaced security & AI picks — local agents, OSINT tooling, and the automation layer most people miss.

Found one of these useful? Forward this to one person who'd get it. That's how the Vault grows.

جرب ولو أداة واحدة هالأسبوع — هذا كل الهدف. أدوات حقيقية تشغلها فعلاً، مش محتوى تتصفحه بس.
عجبك وحدة منها؟ ابعتها لشخص واحد يستفيد. هكذا تكبر الخزينة.



The Toolkit Vault • curated by Hamzah Al-Ramli • goodnbad.info • macOS build